

Grammatical Error Correction on C4_200M

Bag-of-Words → LSTM → Attention → Transformer

Abdullah Ahmed (202200206) · Ibrahim Hanafy (202200518)

CIE 555 — Neural Networks and Deep Learning

Instructor: Dr. Ibrahim Swelam · TAs: Aya Abdelaziz, Mahmoud Farahat

The Problem: Grammatical Error Correction

Rewrite an ungrammatical sentence into its corrected form.

Input (corrupted): "She go to school yesterday and learning many thing."

Output (corrected): "She went to school yesterday and learned many things."

- Errors: tense, agreement, articles, plurals, punctuation, spelling.
- Output is a **whole new sentence** — a sequence-to-sequence task.

The Data: C4_200M Subset

Clean web text automatically corrupted into (corrupted, clean) pairs.

Split	Pairs
Train	850,000
Validation	50,000
Test	100,000
Total	1,000,000

- 16k Byte-Level BPE tokenizer.
- Max length 64 tokens.
- Same splits + tokenizer for every model — no leakage.

Four Models of Increasing Power

#	Model	What it adds
1	Bag-of-Words (TF-IDF)	Baseline — only detects errors
2	LSTM encoder–decoder	First corrector; fixed-size bottleneck
3	LSTM + Bahdanau attention	Decoder re-queries the full source
4	Transformer	Self-attention; no recurrence

Model 1 **detects**. Models 2–4 **correct**.

Model 1 — Bag-of-Words Baseline

- Each pair → two samples: corrupted = ungrammatical (0), clean = grammatical (1).
- Features: TF-IDF, unigrams + bigrams, 50k vocab (fit on train only).
- Classifiers: Logistic Regression & Linear SVM.

Result: ~**62%** detection accuracy — barely above 50% chance. Surface word statistics cannot solve GEC → need seq2seq.

Model 2 — LSTM Encoder–Decoder

- BiLSTM encoder compresses the input into **one** 512-dim state.
- LSTM decoder generates the correction token by token.
- Teacher forcing 1.0→0.5; 19.55M parameters; 3 epochs.
- Test GLEU **0.213** (greedy) — the weakest corrector.

The bottleneck: after the first step the decoder never "sees" the source again — everything is squeezed into 512 numbers.

Model 3 — LSTM + Bahdanau Attention

Keep **all** encoder states; the decoder attends to them every step.

$$\text{score}(s, h) = v \cdot \tanh(W \cdot s + W' \cdot h) \rightarrow \alpha = \text{softmax}(\text{score}) \rightarrow \text{context} = \sum \alpha_i \cdot h_i$$

- Removes the bottleneck; alignment weights are interpretable.
- Padding masked with $-\infty$ before softmax. 29.06M parameters.

The decisive jump: GLEU 0.213 → **0.599** (+0.386). Attention matters most.

Model 4 — Transformer

- No recurrence — self-attention only.
- d_model 256, 8 heads, 3+3 layers, FFN 512.
- Sinusoidal positional encoding; causal mask in the decoder.
- Label smoothing 0.1; warmup learning rate; weight tying.


Smallest & strongest: only **8.05M** parameters (3.6× fewer than Model 3). Best validation GLEU **0.613**.

Comparative Results — Test Set


Model	Decoding	Params	EM	GLEU
LSTM (no attention)	greedy	19.55M	≈0.008	0.213
LSTM (no attention)	beam-4	19.55M	—	0.256
LSTM + Bahdanau	greedy	29.06M	0.024	0.599
LSTM + Bahdanau	beam-4	29.06M	0.026	0.618
Transformer	greedy	8.05M	0.022	0.618
Transformer	beam-4	8.05M	0.020	0.618

- Beam search helps the LSTM, not the (well-calibrated) Transformer.
- Transformer ties Model 3 on GLEU with 3.6× fewer parameters.

Attention Is Interpretable




Bahdanau decoder alignment



Transformer encoder self-attention

Bright cells = which source token the model focuses on.

Behaviour by Sentence Length



All models do best on short sentences; the bottleneck LSTM degrades fastest as sentences get longer.

Grammar-Rule Analysis of Outputs

Learned well (local syntax):

- Subject–verb agreement: "What are this job" → "What is this job"
- Articles: "Why new record" → "Why the new record"
- Verb tense, punctuation, spacing.

Still hard:

- Long-range / semantic edits; rare-word spelling (hallucination).
- Dominant failure = **under-correction** — copying the source unchanged.

The Final Challenge — Dataset Quality

Why does **every** model score only ~2% exact match?

Reference quality is the ceiling — not model capacity

1. **References are paraphrases:** "informed" → "suggest", "startups" → "start-ups" — not grammar errors.
2. **References can be noisy:** target "no brain worky" — not a word.
3. **Many valid corrections exist;** exact match credits exactly one.

→ Report GLEU / F0.5, not EM. Cleaner data helps more than a bigger model.

Conclusion

- **BoW** detects grammaticality at ~62% — cannot correct.
- **LSTM** corrects but is limited by the bottleneck (GLEU 0.213).
- **Bahdanau attention** is the decisive jump (GLEU 0.599).
- **Transformer** ties on GLEU (0.618) with 3.6× fewer params, best validation GLEU (0.613).

Key insight: the binding constraint is **dataset quality**, not architecture — synthetic C4_200M references are often paraphrases or noise.

Thank You

Questions?